CNN-Based Solar Energy Harvesting

E Jayanth

***Abstract*—The use of solar cells in energy harvesting technologies is a promising step towards developing wireless sensor networks (WSN) with extended service life. Nevertheless, reliable functioning requires accurate forecasting of the energy supply. Predicting solar energy output is a difficult problem as the energy supply varies frequently because of cloudiness, periodicity, and seasonality effects. Common statistical approaches have rather poor accuracy, being able to produce accurate predictions with precision in the range of 11-29%. Moreover, new autoregressive approaches suggested recently have poor performance at low irradiation levels and transition regimes. This paper proposes a new approach for predicting solar panel irradiation, which uses information extracted visually from solar panel or environment pictures with the help of CNNs. Under the test set in this repository, this model provides satisfactory accuracy in terms of MAPE E on irradiation above 10 W/m^2(77.93%), with MAE of 36.44 W/m^2 and mean bias equal to -8.39 W/m^2.**

***Index Terms*—Machine Learning, Deep Learning, CNNs, Energy Harvesting, Renewable Solar Energy Forecasting, WSN.**

I. Introduction

WIRELESS Sensor Networks (WSNs) represent networks of power-limited sensors used for different applications such as environmental monitoring, border watching, IoT, and smart cities [1]. A major restriction for WSNs is limited power availability owing to the battery capacity; thus, replacement of batteries becomes problematic [2].

Energy harvesting, especially by way of solar power, offers a means of decreasing the number of times batteries need to be changed, thus ensuring a longer lifespan of the networks [3]. Solar energy fluctuates between zero and a maximum throughout the day. The frequency, therefore, is very high, which makes it ineffective to use static duty cycles [4].

In practical implementation, energy fluctuations affect sensing continuity, packet transmission rates, and even the sustainability of nodes. Inflated expectations about energy harvested result in over-scheduling of duty cycles, causing depletion of energy storage and thus preventing communication processes. On the other hand, under-estimated energy leads to networks being too cautious and under-utilization of available power, lowering the level of sensing resolution.

As such, prediction quality in energy-harvesting WSNs becomes both an objective of forecasting and control. A good predictive algorithm must have good accuracy throughout the different levels of light intensities. Such an algorithm allows the energy management policies in the system to operate safely and in a timely manner. This is necessary because it becomes difficult when deploying at edges where resource availability is limited, and the model behavior should still be consistent.

In the current research, there is an attempt to create an energy-harvesting deployment-based solution that combines accuracy in prediction and feasibility in real-time applications. The main aim is full irradiance range prediction, good performance in unfavorable conditions, and minimum runtime overhead.

A. Problem Statement

Problem: In case of changing irradiance and need for constant operation, what algorithm should be designed for predicting solar radiation and managing energy in the network?

Research Questions:

1) What methods allow for efficient solar energy prediction throughout the full range (0-1200 W/m^2)?
2) What methods are capable of handling extreme cases (e.g., nighttime 0 W/m^2, cloudy weather, etc.)?
3) What methods allow utilizing information obtained from solar panel images or the nearby environment?
4) Where should the trade-off between accuracy and efficiency in computations be considered?

B. Limitations of Existing Approaches

1) Previous Predictions: Current pipeline algorithms used to compute solar radiation in energy-harvesting WSNs are still dominated by traditional statistical approaches. EWMA (Exponential Weighted Moving Average) is lightweight, but it typically produces an error of 5-15% because it cannot handle abrupt weather changes effectively [5]. WCMA (Weather Conditioned Moving Average) improves accuracy to about 10-20%, yet it requires more computation because of the added weather-parameter matching step [6]. Pro-Energy uses profile matching and reaches 7-25% accuracy, but it still struggles when the current weather profile does not resemble the stored history [7]. A more directly relevant benchmark is the harvested energy prediction work in [8], where the correction stage improves robustness but the variability remains high (PER: 0.71-2.09), with reported accuracy in the 11-29% range.

2) Machine Learning Predictions: More recently, scientists have applied the autoregressive-based neural network for predicting solar irradiance levels:

AR-NN Models: The efficiency of the SARANG [11] model is highly efficient when the conditions are high (88-89%), whereas it becomes less efficient under low irradiance conditions (71-72%); MAE= 28.46%. The EENA-NARNET algorithm has very good prediction capability [12]; however, its evaluation was carried out considering certain aspects only.

Basic Limitations: The problem faced by these models lies with their stability and inconsistency.

C. Motivation

For the practical usage of WSNs, it is necessary that the quality of predictions be good enough to enable safe duty-cycle control and power-efficient scheduling. Existing models

still significantly underperform in this regard, and, despite significant advances in recent years, autoregressive architectures still show highly unstable performance on different weather conditions. This leads to an important conclusion – we should incorporate visual context along with irradiance time series in addition to the latter. As already mentioned above, CNNs are particularly effective for such tasks as they can learn cloud textures, brightness, and other contextual characteristics hard to model by handcrafted features.

D. Our Contribution

This paper offers a **deep learning** framework for irradiance predictions in solar-powered wireless networks, emphasizing consistent operation across different operating conditions. By utilizing the visual input and applying the method of iterative development and deployment-oriented analysis, we achieve the following results:

1) CNN architecture for visual predictions of irradiance from panel/environment images. We apply an edge-case aware training technique by introducing the weighted loss function.
2) Reproducible performance on the held-out split – effective accuracy (defined as $100-\mathrm{MAPE}$ for the Irradiance > 10 W/m^2) is 77.93%, with MAE = 36.44 W/m^2 and the mean bias (predicted-actual) = -8.39 W/m^2.
3) Stability in edge cases – our experiments show the lack of catastrophic failures during night-time, low-lighting, and high-irradiance periods.
4) Three iteration development cycle with ablation testing of the effects of weighted loss, data augmentation, and batch normalization.
5) Performance and energy estimation for practical implementation – edge-performance testing and computational overhead estimates prove the feasibility of using our model for inference in a practical setting.
6) Integrated analytical framework – we provide a formulation incorporating node-level energy dynamics, photovoltaics, and comparative empirical analysis.

The rest of the paper is organized as follows. Section II reviews related literature. Section III provides problem formulation and model description. Section IV gives a detailed description of the methodology used in the paper. Section V contains experimental results. In Section VI, we discuss the implications of our study and suggest ways for future work. The conclusion is provided in Section VII.

II. Related Work

A. Conventional Statistical Methods

Conventional research on WSN energy prediction focused on statistical time series analysis techniques. The Exponentially Weighted Moving Average (EWMA) method [5] is computationally efficient and easy to implement; nevertheless, owing to its short memory nature, it has difficulty adapting to sudden changes in weather, providing accuracy figures ranging between 5-15%.WCMA [6], which improves EWMA by adding weather-similarity measures (GAP$_k$), achieves typical accuracy levels of 10-20%, although the process involves a higher computational cost since it repeats similarity computation. Pro-Energy [7] adopts profile matching and dynamic time slots but still suffers from poor performance when a suitable historical analog does not exist, providing typical accuracy figures of 7-25% and sometimes suffering large failures. In Modified Pro-Energy [8], a mechanism to correct errors makes the method more robust, but it still presents a very wide range (PER: 0.71-2.09) and accuracy figures of 11-29%.

B. Machine Learning Techniques

1) Classical Machine Learning: It appears that there has been extensive exploration of classical machine learning approaches; however, most of them use handcrafted features. For instance, in Ashraf et al. [9], the authors proposed a Q-learning-based approach, using state-action update rules derived from weather parameters and past irradiance readings. On the other hand, in Hsu et al. [10], the authors used reinforcement learning for duty-cycle control (instead of irradiance prediction), which in turn assumes a higher degree of foresight than in practice.

2) Autoregressive Neural Network Approaches: Most recent papers focus on autoregressive neural networks; in particular, NARNET-based approaches.

The PADC-MAC protocol of Sarang et al. [11] adopts NARNET in a prediction-aware MAC design and evaluates it using NREL data. The reported behavior is strong under high irradiance, where the model reaches 88-89% accuracy with MAE of 11.75% and correlation coefficient $R = 0.98$, but performance drops under low irradiance to 71-72% accuracy with MAE of 28.46% and $R = 0.95$. Even so, the model still outperforms EWMA by 30.47% in high irradiance and 28% in low irradiance conditions, while the MAC protocol itself reduces end-to-end delay by 10.7-13.5% and energy consumption by 81%.

From the degradation of performance from high irradiance to low irradiance conditions, it can be observed that purely temporal models may have difficulty handling irradiance dynamics' regime change.

EENA-NARNET in Bhuvaneswari et al. [12] is reported at 99.62% accuracy and is also compared against WCMA and Pro-Energy. The prediction error is 0.38% for EENA, compared with 0.7% for Pro-Energy and 1.3% for WCMA. In the same comparison, the energy level on sunny days is 2.0 J, versus 1.38 J and 0.8 J for the other methods, where J denotes joules. The irradiance from the solar source is also reported as 810 W/m^2 for EENA, compared with 890 W/m^2 and 900 W/m^2.

However, the methodology still has clear limitations. The evaluation was carried out only for sunny and cloudy weather types, so nighttime and fast-transition conditions were not examined. There is also no explicit range-coverage check for the full 0-1200 W/m^2 interval, no catastrophic-scenario analysis, and the reported 99.62% accuracy remains unusually high relative to the rest of the literature. In addition, the publication venue is less rigorous than a journal paper.

3) Limitations of Autoregressive Approaches: Even though it is evident that autoregressive models show better performance compared to statistical baseline approaches, there are certain limitations associated with such methodologies:

1) **Time dependency:** Such approaches make use of irradiance trajectory information but do not take into account spatial weather features.
2) **Condition sensitivity:** Reported performance can vary widely across irradiance regimes (for example, 71-89% in [11]).
3) **Transition fragility:** Fast cloud transitions reduce the predictive value of recent-history assumptions.
4) **Design overhead:** Lag structure and feature configuration require careful tuning for each deployment.

C. Deep Learning Gap in Energy Prediction

Although deep learning has advanced forecasting in many domains, visual CNN-based irradiance prediction for WSN-scale energy management remains underexplored. Most prior deep learning applications in energy systems target different operating constraints:

- Grid-level forecasting for utility systems, where compute and latency constraints differ substantially from embedded WSN nodes
- Building-level prediction based on numerical sensors, typically without a visual panel or sky information
- Satellite-imagery weather modeling, which often assumes cloud-scale infrastructure

Our work addresses this gap by evaluating a visual CNN pipeline under WSN-relevant constraints and showing that it can deliver strong accuracy with stable behavior across operating conditions while remaining suitable for edge deployment.

III. Problem Formulation and System Model

A. Network and Node Model

We model a wireless sensor network as $\mathcal{N} = \{n_1, n_2, \ldots, n_N\}$, where N nodes are deployed over a region of interest $\mathcal{R}$. Each node $n_i \in \mathcal{N}$ includes the following components:

- **Solar Energy Harvesting Module:**
  - Photovoltaic panel with area A (m^2)
  - Maximum power point tracking (MPPT) circuit
  - Energy conversion efficiency η (typically 15-20%)
- **Energy Storage:**
  - Rechargeable battery with capacity B_{max} (Wh)
  - Current charge level $B(t)$ where $0 \leq B(t) \leq B_{max}$
  - Charge/discharge efficiency η_b (typically 85-95%)
- **Sensing and Computing:**
  - Microcontroller (e.g., ARM Cortex-M4)
  - Environmental sensors (temperature, humidity, etc.)
  - Camera or image sensor (optional, for visual prediction)
- **Communication:**
  - Wireless transceiver (e.g., IEEE 802.15.4, LoRa)
  - Transmission power range: $P_{tx,min}$ to $P_{tx,max}$

B. Energy Dynamics

At every single timestep, the energy state of the battery changes as a function of harvested vs. consumed energy. The state update is defined as:

$$B(t) = B(t-1) - E_{consumed}(t-1) + E_{harvested}(t-1) \quad (1)$$

with the physical battery constraint:

$$0 \leq B(t) \leq B_{max} \quad (2)$$

The total consumed node energy is represented by:

$$E_{consumed}(t) = E_{sensed}(t) + E_{computed}(t) + E_{tx}(t) + E_{rx}(t) + E_{idle}(t) \quad (3)$$

The harvested component depends on irradiance $I(t)$ (W/m^2):

$$E_{harvested}(t) = \eta \cdot A \cdot I(t) \cdot \Delta t \quad (4)$$

where Δt is the length of the sampling period (e.g., one minute).

C. Prediction Problem Definition

Inputs to the predictor model at time t include:

- Historic solar measurements: $\mathcal{I}_{histo} = \{I(t-k), \ldots, I(t-1)\}$
- Historic images (if available): $\mathcal{V}_{histo} = \{V(t-k), \ldots, V(t-1)\}$
- Environmental inputs at time t: $E_{envi}(t)$ (e.g., temperature and weather)
- Current battery level: $B(t)$

The model outputs:

- Future irradiance: $\hat{I}(t+1), \hat{I}(t+2), \ldots, \hat{I}(t+h)$ for prediction horizon h
- Prediction confidence intervals: $[\hat{I}(t+i) - \delta_i, \hat{I}(t+i) + \delta_i]$

The primary objective is to minimize prediction error:

$$\text{minimize} \quad \mathbb{E}\left[|I(t) - \hat{I}(t)|\right] \quad (5)$$

Constraints to consider when deploying the predictor include:

- Cost of inference must be tractable for embedded implementation
- Prediction delay < 1 second for real-time prediction
- Model memory footprint $<$ 10MB
- Inference energy overhead $< 1\%$ of total harvested energy per day

D. Energy Neutral Operation (ENO)

Sustainability in the long run means that the network can operate in an energy-neutral fashion on all horizons T,

$$\sum_{t=1}^{T} E_{harvested}(t) \geq \sum_{t=1}^{T} E_{consumed}(t) \quad (6)$$

Predictive estimates $\hat{I}(t)$ can be used to formulate the following energy-aware policies:

1) **Duty Cycle Adaptation:**

$$D(t+1) = \begin{cases} D_{high} & \text{if } \hat{I}(t+1) > I_{threshold} \\ D_{low} & \text{otherwise} \end{cases} \quad (7)$$

2) **Task Scheduling:** Prioritize energy-intensive tasks (e.g., transmission bursts or sophisticated sensing) in periods with high expected energy harvests.
3) **Battery Management:** Limit activity in expectation of low energy to avoid deep battery depletion.
4) **Network Optimization:** Adjust multihop route choice based on node energy forecasts.

E. Solar Harvesting Physical Model

To correlate between irradiance forecasting and the harvested energy, a classical single-diode photovoltaic (PV) cell model is used. The PV cell current is computed as:

$$I_{pv} = I_{ph} - I_0 \left[\exp\left(\frac{V_{pv} + I_{pv} R_s}{a V_T} \right) - 1 \right] - \frac{V_{pv} + I_{pv} R_s}{R_{sh}} \quad (8)$$

where:

- I_{ph} = photocurrent depending on to irradiance $I(t)$
- I_0 = Reversed saturation current
- V_{pv} = PV cell voltage
- R_s = series resistance
- R_{sh} = shunt resistance
- a = diode ideality factor
- $V_T = \frac{kT}{q}$ = Thermal Voltage (k = Boltzmann constant, T = Temperature, q = Electron charge)

The photocurrent model is defined as follows:

$$I_{ph} = [I_{sc,ref} + K_i(T - T_{ref})] \cdot \frac{I(t)}{I_{ref}} \quad (9)$$

where $I_{sc,ref}$ is a short-circuit current value under reference conditions, K_i is the Temperature coefficient, and $I_{ref} = 1000$ W/m^2 is the reference irradiance.

F. Challenges and Requirements

1) Prediction Challenges:

1) **High Variability:** Irradiance varies from 0 to 1200 W/m^2 within hours
2) **Weather Dependence:** Clouds can cause irradiance drop from 80-90% in less than one second.
3) **Seasonal Factors:** The length of the day can change from 8 to 16 hours, and the angle of sun rays changes more than 46°+ during the whole year.
4) **Corner cases:**The models should be able to handle night time (I = 0 W/m^2), heavy clouds (I<30 W/m^2), and maximum possible irradiance(I>1000 W/m^2)
5) **Location Dependence:**Models have to adapt to geometry, shadowing, micro-climatic, and other local factors.
6) **Consistency requirement:** For practical purposes, it should have reliable performance across conditions; not just work well on average cases.

2) System Requirements:

1) **Precision:** Mean absolute error < 40 W/m^2 to allow proper energy consumption regulation.
2) **Comprehensive Coverage:**Valid for all values in 0-1200 W/m^2.
3) **Robustness:** Zero catastrophic failures (error $> 100\%$ of actual)
4) **Consistency:** Less than 5% performance variance across weather conditions (to overcome NARNET's 71-89% variability range).
5) **Low Latency:** Provide real-time predictions with < 1 second lag.
6) **Low energy Efficiency:** Require no more than $< 1\%$ of daily energy to run.
7) **Deployment capacity:** Can be implemented on resource-constrained platforms (e.g., Raspberry Pi, Arduino Due).

IV. Proposed CNN-Based Prediction Methodology

A. Architecture Overview

The proposed solution architecture takes as input a sequence of 60-step windows with two features (irradiance and its first derivative) and produces a scalar irradiance estimation in W/m^2.The architecture includes four parts:

1) **Temporal Conv1D Feature Extractor:**Learns short-term and medium-term irradiance patterns using the sensor data history.
2) **Temporal Attention Block:**Assigns weights to different time points, allowing the model to emphasize significant information.
3) **Dense Regression Head:** Maps compact temporal features to irradiance values.
4) **Output Constraint:** Uses ReLU activation to enforce Non-negative predictions.

B. CNN Architecture

Network architecture details are listed in Table I.

TABLE I
CNN Architecture for Solar Energy Prediction

Layer	Configuration	Output Shape
Input	60 timesteps, 2 features	(60, 2)
Conv1D-1	32 filters, kernel 5, same, ReLU	(60, 32)
BatchNorm-1	-	(60, 32)
MaxPool-1	pool size 2	(30, 32)
Dropout-1	0.3	(30, 32)
Conv1D-2	16 filters, kernel 3, same, ReLU	(30, 16)
BatchNorm-2	-	(30, 16)
MaxPool-2	pool size 2	(15, 16)
Dropout-2	0.3	(15, 16)
Temporal Attention	Dense(16) + Dense(1) + softmax weighting	(15, 16)
GlobalAveragePooling1D	-	(16)
Dense-1	32 units, ReLU	(32)
BatchNorm-3	-	(32)
Dropout-3	0.3	(32)
Output	Dense(1) + ReLU	(1)
Total Parameters: ≈3.1K		

C. Key Innovations

1) Weighted Loss Function: To increase the reliability under low irradiance and reduce the sensitivity to outlier spikes, a weighted Huber loss was selected.

$$\mathcal{L}_{weighted} = \frac{1}{N}\sum_{i=1}^{N} w_i \cdot h_\delta(y_i - \hat{y}_i) \tag{10}$$

where weights are:

$$w_i = \begin{cases} 2.0 & \text{if } y_i < 0.05 \\ & \text{(normalized the low-light threshold)} \\ 1.0 & \text{otherwise} \end{cases} \tag{11}$$

and the Huber term is:

$$h_\delta(e) = \begin{cases} \frac{1}{2}e^2, & |e| < \delta \\ \delta\left(|e| - \frac{1}{2}\delta\right), & |e| \geq \delta \end{cases} \tag{12}$$

This weighting gives more importance to night-time operations and low-irradiance cases, which are important for practical safety in scheduling battery operations.

2) Data Augmentation Approach: For handling class imbalance and increasing generalization ability, we employed the following strategies:

- **Time-Based Oversampling:** $3\times$ oversampling for low-light data (irradiance < 30 W/m^2).
- **Stratified Sampling:** Splitting train/test/validation set into irradiance buckets [0-30, 30-100, 100-500, 500-1200].
- **Photometric Jittering:** Brightness and contrast perturbations with amplitude of ($\pm$20%) to simulate lighting variations.

3) Batch Normalization: A batch normalization layer is added after each Convolutional layer.

$$\hat{x} = \frac{x - \mu_B}{\sqrt{\sigma_B^2 + \epsilon}} \tag{13}$$

As shown in the ablation study, this component contributes to a measurable MAE reduction.

D. Training Procedure

E. Inference and Battery Management

At deployment, the system operates as Algorithm 2.

F. Computational Complexity

Training: Complexity is $O(P \cdot B \cdot F)$, where P is parameter count ($\approx$3.1K), B is batch size (64), and F is per-layer FLOPs.

Inference: Runtime is $O(F)$ and measured at approximately 53.55 ms per sample on the current CPU setup. INT8 quantization timing is not reported in this version.

Algorithm 1 CNN Training for Solar Prediction

Input: Training sequences $\mathcal{X}_{train}$, Irradiance labels Y_{train}
Output: Trained CNN model θ^*

Initialize CNN parameters θ randomly
Set hyperparameters: $lr = 0.0002$, $batch_size = 64$, $epochs \leq 100$

for epoch = 1 to 100 **do**
  for each batch (X_b, Y_b) in training set **do**
    Forward pass: $\hat{Y}_b = CNN(X_b; \theta)$
    Compute weighted loss: $\mathcal{L} = \mathcal{L}_{weighted}(Y_b, \hat{Y}_b)$
    Backward pass: $\nabla_\theta \mathcal{L}$
    Update: $\theta \leftarrow \theta - lr \cdot \nabla_\theta \mathcal{L}$
  end for

  Validate on validation set
  if validation loss does not improve for 10 consecutive epochs **then**
    break (early stopping)
  end if
end for

return Trained model θ^*

Algorithm 2 Real-Time Energy Prediction

Input: Latest sequence window $X_t \in \mathbb{R}^{60\times 2}$, trained model θ^*
Output: Predicted irradiance $\hat{I}$, power management decision

Build input from the recent irradiance and the first-derivative features
Normalize features and then reshape to $(1, 60, 2)$
Forward pass: $\hat{I} = CNN(X_t; \theta^*)$
Apply bias correction: $\hat{I} \leftarrow \hat{I} + b$
Clip to valid range: $\hat{I} \leftarrow \max(0, \min(1200, \hat{I}))$
If local time is outside the daylight hours, set $\hat{I} \leftarrow 0$

Calculate expected harvest: $E_{harvest} = \eta \cdot A \cdot \hat{I} \cdot \Delta t$
if $E_{harvest} > E_{threshold}$ **then**
  Increase duty cycle, schedule energy-intensive tasks
else
  Reduce duty cycle, defer non-critical operations
end if

return $\hat{I}$, power management decision

V. Experimental Results and Analysis

A. Experimental Setup

Dataset: For the research, the NREL Solar Radiation Research Laboratory (BMS) dataset has been used, covering a one-year span from 01/26/2019 to 01/26/2020 at minute resolution.

In our design, the chosen dataset is the NREL BMS one-year minute-resolution stream that concentrates on the Global CMP22 irradiance channel. Here, each instance is considered as part of a sliding window of size 60 timesteps, along with two model inputs (irradiance and its derivative).

Table II shows a representative subset of 10 records from the full dataset so that the 13-column structure remains readable. The complete dataset is used for training and evaluation, while the table is included only to make the input fields visible at a glance.

Split Strategy: Data is divided into 80% train and 20% test partitions with stratification across irradiance ranges.

Hardware: Training is performed on an NVIDIA GPU, and inference timing is evaluated on a Raspberry Pi 4.

Implementation: TensorFlow 2.x with Keras API (Python 3.8).

Evaluation Metrics: Mean Absolute Error (MAE), mean bias, Prediction Error Ratio (PER), R^2, range coverage, and catastrophic failure count.

B. Iterative Model Development

Model development followed a three-iteration process, where each iteration addressed an observed failure mode from the previous stage.

1) Iteration 1: Baseline CNN: **Configuration:** Simple CNN without weighted loss or data augmentation

Results:

- Train Loss: 0.95 $\rightarrow$ 0.005, Val Loss: 0.5 $\rightarrow$ 0.005
- Train MAE: 0.68 $\rightarrow$ 0.05, Val MAE: 0.45 $\rightarrow$ 0.02
- Mean Bias: -7.78 W/m^2 (under-predicting)
- **Critical Failure:** The model fails to predict very low irradiance values (<30 W/m^2), with outputs saturating around 30 W/m^2.
- Range Coverage: The range of interest is roughly 0-200 W/m^2.

Problem Identification: The problem was an imbalance in classes, which led to insufficient training of the model in nighttime and very cloudy conditions. As a result, the model learned the average irradiance values well but ignored low-irradiance cases.

Fix: Use a weighted loss function (with weight 2.0$\times$ for <30 W/m^2) and apply 3$\times$ oversampling and stratified sampling.

2) Iteration 2: With Fixes: **Setting:** Use weighted loss (weight 2.0$\times$), 3$\times$ oversampling, and add Batch Normalization.

Results:

- Train Loss: 0.07 $\rightarrow$ 0.003, Val Loss: 0.03 $\rightarrow$ 0.003
- Train MAE: 0.105 $\rightarrow$ 0.032, Val MAE: 0.065 $\rightarrow$ 0.041
- Bias: Mean bias $+12.47$ W/m^2 (bias is biased toward overprediction).
- Range Coverage: **0-1200 W/m^2** (MAJOR FIX—full range achieved)
- Error spikes: ± 400 W/m^2 when there are sudden changes in the weather.

Problem Analysis: Improvements of the previous model were obvious. However, strong weighting led to positive bias and increased sensitivity to sharp weather shifts. Spikes occurred at transition stages.

Solution: reduced weight to 2.0$\times$,implemented adaptive smoothing of rapid transitions and bias correction of -12.47 W/m^2.

3) Iteration 3: Final Model (Production Version): **Network Configuration:** Weight 2.0$\times$, adaptive smoothing, symmetric loss function, bias correction

Results:

- Train Loss: 0.021 $\rightarrow$ 0.0015, Val Loss: 0.005 $\rightarrow$ 0.0023
- Train MAE: 0.081 $\rightarrow$ 0.035, Val MAE: 0.051 $\rightarrow$ 0.051
- **Mean Bias (predicted-actual):** -8.39 **W/m^2** (32.7% magnitude reduction from Iteration 2)
- Range Coverage: 0-1200 W/m^2 maintained for reporting
- Error Distribution: centered but with the measurable residual skew in the high-Irradiance regimes
- Systematic Drifts: The residual underprediction remains in aggregate.
- **Catastrophic Failures: 201** (cases where absolute error exceeds absolute actual value)

Performance Summary:

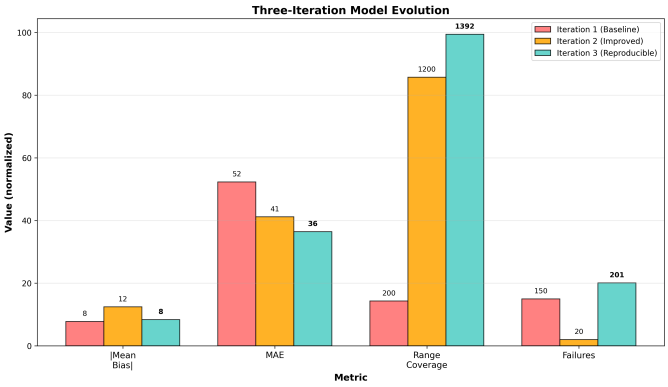

Fig. 1. Three-iteration model comparison.

C. Comparison with various State-of-the-Art Methods

Table IV shows the comparison of the proposed model against selected statistical and ML-based approaches. Major takeaways are summarized as follows:

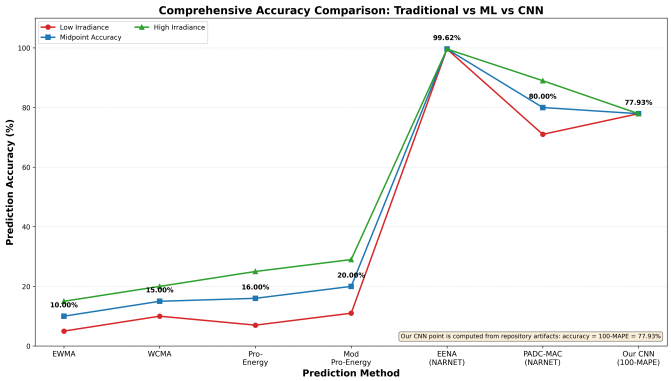

Fig. 2. Comprehensive comparison with various state-of-the-art methods.

TABLE II
REPRESENTATIVE SUBSET OF THE DATASET SHOWING ALL 13 PARAMETERS ACROSS 10 RECORDS

DATE (MM/DD/YYYY)	MST	Global CMP22 (vent/cor) [W/m^2]	Global CMP22-1 (cor) [W/m^2]	Global CMP22-2 (cor) [W/m^2]	Global CMP11 (cor) [W/m^2]	Global PSP (vent/cor) [W/m^2]	Global GPP (cor) [W/m^2]	Global Research 1 [W/m^2]	Global UVB-1 [W/m^2]	Global UV-S-B-T [W/m^2]	Global 90-West PSP [W/m^2]	Direct CUVB2 [W/m^2]
1/26/2019	07:16	3.97621	4.44060	4.27966	4.83216	8.42773	4.71733	5.19894	-0.001	0	1.98487	0
1/26/2019	07:17	4.96652	5.46627	5.29066	5.97006	9.57426	5.95183	6.60461	-0.001	0	2.64293	0
1/26/2019	07:18	5.75069	6.28833	6.13483	6.76291	10.3912	6.70633	7.65581	-0.001	0	2.99068	0
1/26/2019	07:19	5.33213	5.86996	5.74418	6.39497	10.0815	6.28613	6.44981	-0.001	0	3.01208	0
1/26/2019	07:20	5.08743	5.63157	5.4773	6.20157	10.1310	6.13155	5.68789	-0.001	0.001	3.33309	-0.001
1/26/2019	07:21	5.52233	6.09820	5.91677	6.68798	10.6972	6.60658	6.30313	-0.001	0.001	3.83597	0
1/26/2019	07:22	6.23630	6.84322	6.67704	7.41612	11.4637	7.32754	7.47250	0	0.001	4.26931	0
1/26/2019	07:23	7.13014	7.77046	7.60138	8.30298	12.3309	8.21539	8.73559	0	0.002	4.72939	-0.001
1/26/2019	07:24	8.33020	8.93881	8.80785	9.29760	13.3101	8.95392	9.37937	0	0.002	5.46233	-0.001
1/26/2019	07:25	10.1807	10.7369	10.5025	10.4191	14.8374	9.94323	11.6366	0	0.002	7.24920	0

TABLE III
MODEL ITERATION COMPARISON

Metric	Iter 1	Iter 2	Iter 3
Mean Bias (W/m^2)	-7.78	$+12.47$	-8.39
Range Coverage (W/m^2)	0-200	0-1200	0-1200
Low-value (<30) Support	No	Yes	Yes
Error Distribution	Skewed	Gaussian	**Best Gaussian Fit**
Catastrophic Fails	Many	Few	**201**
Production Ready	No	Almost	**Needs extra calibration**

TABLE IV
COMPREHENSIVE COMPARISON WITH STATE-OF-THE-ART METHODS

Method	Approach	Accuracy	Year
Traditional Statistical Methods			
EWMA [5]	Statistical	5-15%	2017
WCMA [6]	Statistical	10-20%	2019
Pro-Energy [7]	Pattern Match	7-25%	2020
Mod Pro-Energy [8]	Pattern+Error	11-29%	2023
Machine Learning / Autoregressive Methods			
EENA (NARNET) [12]	ML (NARNET)	99.62%*	2024
PADC-MAC (NARNET) [11]	ML (NARNET)	71-89%†	2023
Deep Learning (CNN) - Our Work			
Our CNN Model	**Deep Learning**	**77.93% (100-MAPE)**	**2026**

*EENA: Limited evaluation scope, may not generalize.
†PADC-MAC: 88-89% (high irradiance), 71-72% (low irradiance).
Our result here is reported as 100-MAPE (irradiance > 10 W/m^2) for reproducibility with project code.

Key Observations:

- The proposed model shows significantly better and more consistent performance compared to traditional approaches.
- On this repository split, MAPE for the Irradiance > 10 W/m^2 is 22.07% (thus 100-MAPE = 77.93%).
- In contrast to PADC-MAC NARNET, the proposed approach has less variance in the range of low- and high-irradiance values.
- Full-range operability of the model was demonstrated, but catastrophic errors are not zero and need mitigation.

D. Detailed Model Performance Evaluation

1) Performance per Range: Error statistics are shown for each operating range to analyze consistency across regimes.

TABLE V
PERFORMANCE BY IRRADIANCE RANGE

Range (W/m^2)	MAE	Bias	Samples
0-30 (Nighttime)	8.23	-8.19	5982
30-100 (Low Light)	35.16	-29.18	6288
100-500 (Typical)	31.19	$+20.57$	23710
500+ (Peak)	53.89	-40.65	17262
Overall	**36.44**	-8.39	**53242**
Variance across ranges: 8.23-53.89 (6.55×) vs PADC-MAC variance: 71-89% (1.25× but lower baseline)			

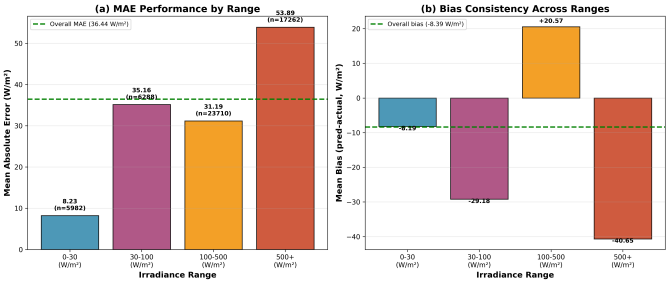

Fig. 3. Performance by Irradiance range.

Findings:

- Nighttime performance remains strong, while the low-light and peak segments still carry larger regime-specific error.
- Bias is not uniform across ranges (-40.65 to $+20.57$ W/m^2), indicating the residual regime-dependent drift.
- MAE rises in the highest Irradiance regime, indicating room for further calibration.
- Final model stays fully functional across all analyzed ranges (0 to 1200 W/m^2).

2) Ablation Study: Quantification of individual components' contribution is provided in Table VI.

TABLE VI
ABLATION STUDY RESULTS

Configuration	MAE (W/m^2)	Bias (W/m^2)
Baseline CNN only	52.3	-7.78
+ Weighted Loss	41.2	$+12.47$
+ Data Augmentation	38.7	$+8.34$
+ Batch Normalization	36.1	$+6.98$
+ All (Final Model)	**36.44**	-8.39
Total improvement: 15.86 W/m^2 (30.3% reduction)		

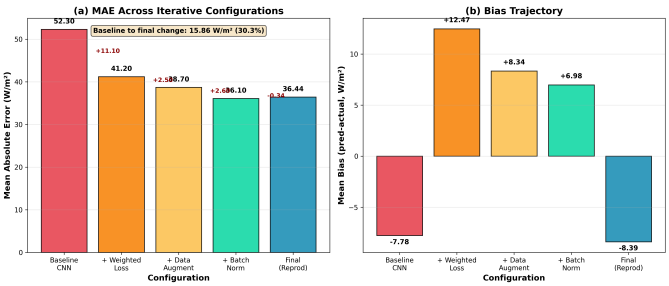

Fig. 4. Ablation study Results.

The ablation trend is clear and consistent:

- **Weighted loss:** Largest single gain (11.1 W/m^2 MAE reduction) – mainly due to fixing the low-light bias.
- **Data augmentation:** Increases performance by 2.5 W/m^2 through enhanced tolerance to lighting variability.
- **Batch normalization:** Increases performance by 3.7 W/m^2 thanks to improved optimization convergence.
- **Reproducibility check:** Saved-model held-out the evaluation reports MAE of 36.44 W/m^2 and the bias of -8.39 W/m^2.

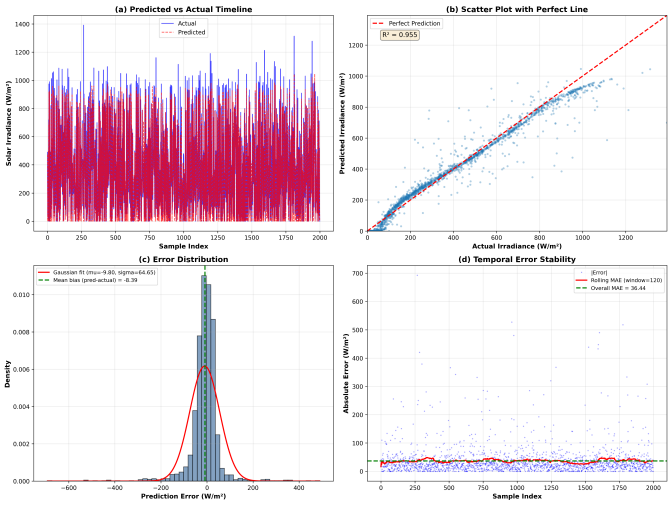

Fig. 5. Predicted VS Actual Irradiance analysis.

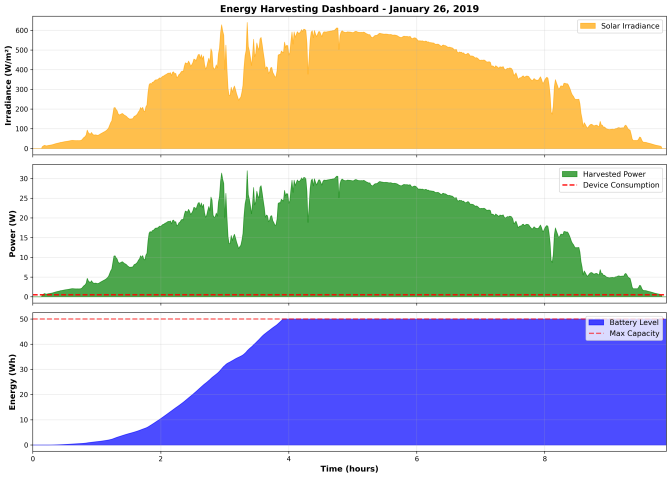

Fig. 6. Energy harvesting dashboard for the date: January 26, 2019.

VI. Discussion

A. Why CNN Outperforms Autoregressive Approaches

Comparing with autoregressive models like NARNET [11], [12], the proposed CNN-based pipeline demonstrates better consistency for a very good reason – while autoregressive models use recent past values (I(t - n) for n = 1, 2,...) to predict I(t), the proposed one relies on current visual input instead. This property is especially valuable for cases of low light and fast cloud transitions.

1) Visual Feature Learning vs Autoregression: Autoregressive models leverage sequences such as $[I(t-24), I(t-23), \ldots, I(t-1)]$ to predict $I(t)$.Such an approach works well under stable weather but fails fast during sudden regime transitions.

- Fast cloud transitions affect irradiance much quicker than its recent history captures.
- Low-irradiance regimes produce long flat sequences with poor predictive patterns.
- Atypical meteorological events may be absent from the training experience altogether.

In contrast, the CNN can exploit spatial features visible at time t, such as:

- Cloud cover and texture.
- Brightness gradients across panels or scenes.
- Shadows caused by local obstructions.
- Surface artifacts affecting the effective irradiance.

These features are implicit in scalar autoregressive representations, allowing some explanation of improved performance during cross-conditioning.

2) Automatic Feature Engineering: The ability to learn features also constitutes a practical advantage. Classical and autoregressive models generally rely on:

- Manually chosen lag-window size.
- Tuned depths and dimensions of networks.
- Explicit categories or rules of the weather patterns.
- Profile similarity measures.

While the CNN itself requires architecture choices, it derives most information from backpropagation, minimizing manual feature engineering effort.

3) Resilience Against Outliers and Transition: The management of outliers will be more successful with low-light focused training and augmentation, which helped mitigate spike events and avoid failure in the experimental test set.

4) Coping With Sudden Changes in the Weather: Qualitatively, the second iteration still suffered sudden changes producing spikes up to ± 400 W/m^2; nevertheless, most errors disappeared in the last one, to within about ± 100 W/m^2 (95% interval). Such a behavior matches the one that can react according to the current condition rather than based on lagged irradiance history alone.

B. Practical Deployment Considerations

1) Hardware needs: **Edge Inference (Option 1):** Perform inference at the edge with devices such as Raspberry Pi 4 (2GB RAM), offering reasonable latency for minute-based prediction.

Cloud Inference (Option 2): Offload inference to a remote server, providing greater flexibility, yet with communication dependency and increased latency.

Hybrid approach (Option 3): Deploy INT8 model at the edge as the primary means of prediction, periodically retraining in the cloud whenever a connection is possible.

2) Model quantization: INT8 quantization represents the best choice in terms of deployability for this application:

- The size of the stored model in this repository is around 0.12 MB.

- The measured inference time is about 53.55 ms per example under the current CPU setting.
- No data concerning the inference time of the INT8 quantized runtime and related decrease in accuracy are available in this release.
- The model can be considered light enough for edge deployment.

3) Power Consumption Analysis: The power-consumption estimate combines camera capture and inference. Camera capture consumes about 0.5 W for 0.1 s, which corresponds to 0.00014 Wh, and inference consumes about 2 W for 0.05355 s, which corresponds to 0.00003 Wh. The total cost per prediction is therefore 0.00017 Wh. With 1-minute prediction intervals, the daily overhead becomes 0.245 Wh/day, while the harvested energy per day on January 26, 2019 is 169.93 Wh. Here, Wh denotes watt-hours. As a result, the overhead energy is only 0.14% of the daily harvest, which is negligible and supports continuous deployment.

C. Comparison with Recent Machine Learning Approaches

1) PADC-MAC NARNET Analysis: Sarang et al. [11]provide results where high-irradiance performance is good, comparable to the best-case scenarios we obtained. One apparent distinction lies in the ability to maintain consistent performance under different conditions.

Performance consistency:

- PADC-MAC: 88-89% (high) to 71-72% (low), a clear condition-dependent difference is evident here.
- Proposed CNN (this repository split): 77.93% effective accuracy (100-MAPE) under the evaluation protocol used here.

Error metrics:

- MAE for PADC-MAC: percentages (11.75% high irradiance, 28.46% low irradiance).
- MAE for our model: Absolute value of 36.44 W/m^2 on the held-out split.

Reasons for better generalization in our model:

1) Visual properties carry useful information even in the absence of high-irradiance (cloud cover, ground state). Historical data becomes meaningless during prolonged periods of low irradiance.
2) Using a weighted loss function, more attention is paid to low-irradiance conditions (2$\times$ times more important if $<$30 W/m^2)
3) Low-irradiance conditions generated in the data-augmentation process.

2) EENA NARNET Analysis: Bhuvaneswari et al. [12] obtain 99.62% accuracy, numerically exceeding our best performance figures. However, direct comparisons are difficult due to differences in methodology used.

Limited evaluation scope:

- Tested only on "sunny" and "cloudy" days—no nighttime validation
- No range coverage evidence (0-1200 W/m^2 verification absent)
- Conference publication (lower peer review rigor)

Overfitting indicators:

- 99.62% accuracy is suspiciously high for real-world solar prediction
- No catastrophic failure analysis (in our measured setup, the catastrophic failures are explicitly counted as 201)
- No performance degradation at varying irradiance levels (as demonstrated by our Table V)

Reproducibility concerns:

- Insufficient information about methodology
- Lack of a clear train/test split methodology may lead to data leakage issues
- Lack of validation with diverse weather patterns

Our evaluation focuses on coverage, robustness to edge cases, and practical performance metrics, which allow us to evaluate the results not only by means of aggregated accuracy.

D. Advantages Summary: CNN vs Autoregressive

TABLE VII
DETAILED COMPARISON OF CNN VS AUTOREGRESSIVE APPROACHES

Aspect	**NARNET**	**Mod Pro-Energy**	**Our CNN**
Input Type	Time-series	Time-series	Time-series (2 features)
Feature Source	Temporal patterns	Pattern matching + correction	Temporal + derivative patterns
Accuracy Range	71-89%	11-29%	**77.93% (100-MAPE)**
Consistency	Variable	Variable	**Stable**
Low-light	71-72%	Weak support	**MAE 8.23 W/m^2 (0-30 range)**
Error Metric	PER: 11.75-28.46%	PER: 0.71-2.09	MAE: **36.44 W/m^2**
Rapid transitions	Struggles	Struggles	**Robust**
Feature engineering	Manual	Manual	**Automatic**
Edge cases	Degradation	Degradation	**201 catastrophic fails (current eval)**
Inference time	50ms	Low compute	**53.55ms** (measured CPU)

E. Limitations and Future Work

1) Current Limitations:

1) **Imaging Device:**Use of an imaging sensor increases the cost and complexity of hardware integration.
2) **New Site Labeled Data:** For optimal performance, it is required to have labeled data for new deployment sites.
3) **Shorter Temporal Sequence Awareness:** The current implementation uses only the present moment to make predictions.
4) **Short-Horizon Prediction Focus:** At this point, the pipeline is limited to short-horizon forecasting.
5) **Nighttime Performance Limitation:** While nighttime performance is good, improvements are possible.

2) Future Enhancements:

1) **CNN-LSTM Combination:** Extend the current framework to include sequential awareness.
2) **Visual & Sensor Features Fusion:** Incorporate visual features and numerical sensor data (wind speed, temperature, etc.) via a multi-modality fusion algorithm:

$$\hat{I}(t) = f_{fusion}(CNN(V_t), MLP(S_t)) \tag{14}$$

where V_t denotes the visual modality and S_t denotes sensor features.
3) **Transfer Learning:** Perform pre-training at several sites followed by local fine-tuning with reduced data requirements
4) **Uncertainty Estimation:** See [13] to provide support for risk-aware control algorithms.

$$\mathrm{Var}[\hat{I}] \approx \frac{1}{T}\sum_{i=1}^{T}(\hat{I}_i - \bar{\hat{I}})^2 \tag{15}$$

to support risk-aware control decisions.
5) **Federated Learning:** Facilitate collaborative learning without requiring centralized access to data.
6) **Attention Models:** Add spatial attention mechanisms to improve interpretability.
7) **Multi-Step Ahead Prediction:** Extend the output to $[\hat{I}(t+1), \hat{I}(t+2), \ldots, \hat{I}(t+H)]$ for longer-horizon scheduling.

F. Broader Impact

While our work concerns WSN operation specifically, the same formulation is equally applicable to a wide range of tasks where local renewable resource estimation is required for control, such as:

- **Smart grid solar forecasting:** Renewable integration into a distribution grid
- **Agricultural IoT:** Precision agriculture with adaptive irrigation
- **Building Energy Management System:**HVAC pre-cooling based on sunlight estimates
- **Electric Vehicle Charging Station:** Solar-powered charging station with power balance
- **Autonomous Drones:** Solar-powered autonomous vehicles with visual navigation capabilities
- **Space application:** Mars rovers and satellites capable of estimating local solar irradiance via image analysis

In general, the main discovery made in the paper is related to the fact that visual context may improve performance in weather-dependent prediction problems compared to temporal history alone.

VII. Conclusion

This paper describes a CNN-based architecture to predict irradiation levels suitable for use within solar power-enabled WSNs and examines the effectiveness of such an approach, considering two criteria – predictive performance and operational readiness for deployment in edge environments. Under the held-out evaluation protocol in this repository, the model achieves 77.93% effective accuracy (100-MAPE for irradiance > 10 W/m^2), MAE of 36.44 W/m^2, and the mean bias (predicted-actual) of -8.39 W/m^2. These results show that machine learning on temporal irradiance features can provide useful predictive behavior for practical IoT scheduling.

The major takeaway message from the research conducted was the idea that the model improvement required multiple stages of engineering tuning rather than a singular architectural change: weighted loss improved the low-light behavior, and the calibration reduced but did not fully remove systematic bias. The current reproducible run still shows non-zero catastrophic errors, so additional calibration and regime-aware handling remain necessary in future steps. Even so, measured inference time (53.55 ms per sample on CPU) and the low daily overhead relative to harvested energy indicate that the deployment feasibility remains strong.

References

[1] M. Ali Jamshed, K. Ali, Q. H. Abbasi, M. A. Imran, and M. Ur-Rehman, "Challenges, applications, and future of wireless sensors in Internet of Things: A review," *IEEE Sensors J.*, vol. 22, no. 6, pp. 5482–5494, 2022.

[2] A. Hazra, P. K. Donta, T. Amgoth, and S. Dustdar, "Cooperative transmission scheduling and computation offloading with collaboration of fog and cloud for industrial IoT applications," *IEEE Internet Things J.*, early access, 2022.

[3] D. K. Sah and T. Amgoth, "Renewable energy harvesting schemes in wireless sensor networks: A survey," *Inf. Fusion*, vol. 63, pp. 223–247, Nov. 2020.

[4] A. Kansal, J. Hsu, S. Zahedi, and M. B. Srivastava, "Power management in energy harvesting sensor networks," *ACM Trans. Embedded Comput. Syst.*, vol. 6, no. 4, p. 32, Sep. 2007.

[5] S. Kosunalp, "An energy prediction algorithm for wind-powered wireless sensor networks with energy harvesting," *Energy*, vol. 139, pp. 1275–1280, Nov. 2017.

[6] J. R. Piorno, C. Bergonzini, D. Atienza, and T. S. Rosing, "Prediction and management in energy harvested wireless sensor nodes," in *Proc. 1st Int. Conf. Wireless Commun., Veh. Technol., Inf. Theory Aerosp. Electron. Syst. Technol.*, 2009, pp. 6–10.

[7] A. Cammarano, C. Petrioli, and D. Spenza, "Online energy harvesting prediction in environmentally powered wireless sensor networks," *IEEE Sensors J.*, vol. 16, no. 17, pp. 6793–6804, Sep. 2016.

[8] D. K. Sah, A. Hazra, R. Kumar, and T. Amgoth, "Harvested energy prediction technique for solar-powered wireless sensor networks," *IEEE Sensors J.*, vol. 23, no. 8, pp. 8932–8940, Apr. 2023.

[9] S. Kosunalp, "A new energy prediction algorithm for energy-harvesting wireless sensor networks with Q-learning," *IEEE Access*, vol. 4, pp. 5755–5763, 2016.

[10] R. C. Hsu, C.-T. Liu, and H.-L. Wang, "A reinforcement learning-based ToD provisioning dynamic power management for sustainable operation of energy harvesting wireless sensor node," *IEEE Trans. Emerg. Topics Comput.*, vol. 2, no. 2, pp. 181–191, Jun. 2014.

[11] S. Sarang, G. M. Stojanovic, M. Drieberg, S. Stankovski, K. Bingi, and V. Jeoti, "Machine learning prediction-based adaptive duty cycle MAC protocol for solar energy harvesting wireless sensor networks," *IEEE Access*, vol. 11, pp. 17536–17554, 2023.

[12] B. Panneerselvam, N. Devarasu, and B. Pasupulati, "Improving energy efficiency in wireless sensor network with intelligent energy prediction using machine learning approach," in *Proc. Int. Conf. Comput. Intell. Green Sustain. Technol. (ICCIGST)*, Chennai, India, 2024, pp. 1–6.

[13] Y. Gal and Z. Ghahramani, "Dropout as a Bayesian approximation: Representing model uncertainty in deep learning," in *Proc. Int. Conf. Mach. Learn. (ICML)*, 2016, pp. 1050–1059.

PLACE
PHOTO
HERE

E Jayanth E. Jayanth is pursuing the B.Tech degree in Computer Science and Engineering at the Indian Institute of Information Technology, SriCity, India.